

431 Class 10

Thomas E. Love

2017-09-28

Today's Agenda

- Forming Project Task B Groups
- The Western Collaborative Group Study
 - Dealing with Factors
 - Building Tables Effectively
 - Dealing with Missingness
 - Scatterplot and Correlation Matrices

15 Questions Dr. Love plans to include in the Survey

The following items will be included in the survey. As a result, you will not want to ask these questions in your Task B, although you should consider these groupings as candidates for application in your research questions.

These 8 items will be provided in groups after the application of cutpoints we will identify together after the survey is complete.

1. In what year were you born?
2. How would you rate your current health overall (Excellent, Very Good, Good, Fair, Poor)
3. For how long, in months, have you lived in Northeast Ohio?
4. What is your height in inches? (If you are five feet, eight inches tall, please write 68 inches. To convert from centimeters to inches, multiply your height in centimeters by 0.3937, and then round the result to the nearest inch.)
5. What is your weight in pounds? (To convert from kilograms to pounds, multiply your weight in kilograms by 2.2046, and then round the result to the nearest pound.)
6. What is your pulse rate, in beats per minute? (Please either use a tracking device, or count your pulse for 15 seconds then multiply by 4)
7. Last week, on how many days did you exercise? (0 - 7)
8. Last night, how many hours of sleep did you get?

The following 7 items will have yes/no responses, and thus produce binary groups for analysis.

1. Were you born in the United States?
2. Is English the language you speak better than any other?
3. Do you identify as female?
4. Do you wear prescription glasses or contact lenses?
5. Before taking 431, had you ever used R before?
6. Are you currently married or in a stable domestic relationship?
7. Have you smoked 100 cigarettes or more in your entire life?

Project Task B groups

We need ten such groups, each with about 5 people involved.

Google Form is available at <https://goo.gl/forms/WaQOdCEAW0wxdjJh2> and needs to be done by 5 PM today.

Task B meetings in class will be held next Tuesday, and also on 2017-10-12.

Details on Task B specified at <https://github.com/thomaseLove/431project>

The Form's Questions...

Fall 2017 Project Task B Groups

This is the form to specify the group name and membership for Task B. Only one person from your group should fill out this form. The Task B groups will be formed in class on 2017-09-26. This form needs to be submitted by noon on 2017-09-27. If you have questions, contact Dr. Love directly.

Your email address (**tel3@case.edu**) will be recorded when you submit this form. Not you? [Switch account](#)

* Required

What is the name of your Task B group? *

100 characters or less, please.

Your answer

Please select the names of the members of your group from the list below.

Your group must include 4-6 people, in total. Be sure to check the box for each group member, including yourself.

In Our Group

Albar, Zainab

☐

Asagba, Oghenerukema

☐

Today's R Setup

```
library(Epi); library(viridis); library(broom)
library(GGally); library(mice)
library(forcats); library(tidyverse)

source("Love-boost.R")
```

Cleaning up Loose Ends in the VHL Study

```
VHL <- read.csv("vonHippel-Lindau.csv") %>% tbl_df
```

Von Hippel - Lindau study Codebook

- p.ne = plasma norepinephrine (pg/ml)
- tumorvol = tumor volume (ml)
- disease = 1 for patients with multiple endocrine neoplasia type 2
- disease = 0 for patients with von Hippel-Lindau disease

We want to add a new variable (factor) called diagnosis, which takes the values von H-L or neoplasia.

Creating a Factor to represent disease diagnosis

We want to add a new variable, specifically a factor, called `diagnosis`, which will take the values `von H-L` or `neoplasia`.

- Recall `disease` is a numeric 1/0 variable (0 = `von H-L`, 1 = `neoplasia`)
- Use `fct_recode` from the `forcats` package...

```
VHL <- VHL %>%  
  mutate(diagnosis = fct_recode(factor(disease),  
                                "neoplasia" = "1",  
                                "von H-L" = "0")  
  )
```


Now, what does VHL look like?


VHL

```
# A tibble: 37 x 5
```

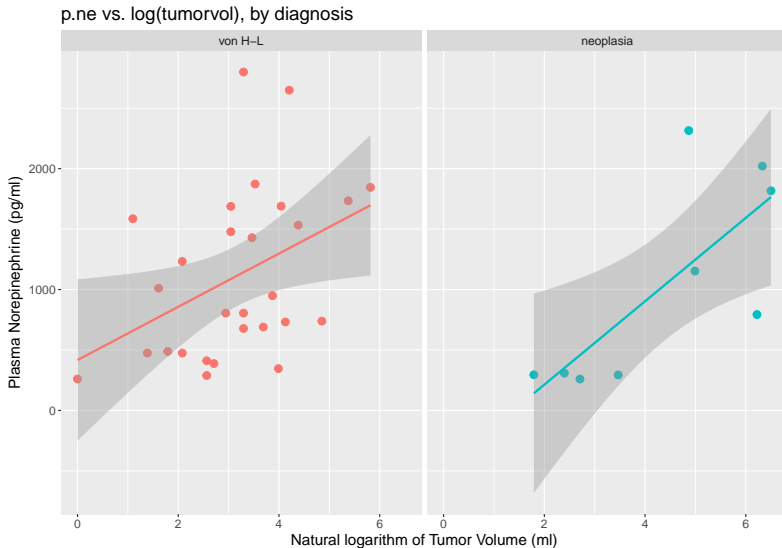
	id	disease	p.ne	tumorvol	diagnosis
	<int>	<int>	<int>	<int>	<fctr>
1	101	0	289	13	von H-L
2	102	1	294	32	neoplasia
3	103	0	2799	27	von H-L
4	104	0	2649	67	von H-L
5	105	0	346	54	von H-L
6	106	0	1690	57	von H-L
7	107	0	805	19	von H-L
8	108	1	1153	147	neoplasia
9	109	0	678	27	von H-L
10	110	1	1817	665	neoplasia

```
# ... with 27 more rows
```

Compare the patients by diagnosis



Facetted Scatterplots by diagnosis



Model accounting for different slopes and intercepts

```
model2 <- lm(p.ne ~ log(tumorvol) * diagnosis, data = VHL)
model2
```

Call:

```
lm(formula = p.ne ~ log(tumorvol) * diagnosis, data = VHL)
```

Coefficients:

```
              (Intercept)
                417.2
        log(tumorvol)
                220.0
diagnosisneoplasia
               -893.3
log(tumorvol):diagnosisneoplasia
                124.8
```

Model 2 results

$$p.ne = 417 + 220 \log(\text{tumorvol}) - 893 (\text{diagnosis} = \text{neoplasia}) + 125 (\text{diagnosis} = \text{neoplasia}) * \log(\text{tumorvol})$$

where the indicator variable $(\text{diagnosis} = \text{neoplasia}) = 1$ for neoplasia subjects, and 0 for other subjects...

- Model for $p.ne$ in von H-L patients:
 - $417 + 220 \log(\text{tumorvol})$
- Model for $p.ne$ in neoplasia patients:
 - $(417 - 893) + (220 + 125) \log(\text{tumorvol})$
 - $-476 + 345 \log(\text{tumorvol})$

Model 2 Predictions

What is the predicted p.ne for a single new subject with tumorvol = 55 ml (so $\log(\text{tumorvol}) = 4.01$) in each diagnosis category?

```
predict(model2, newdata = data_frame(tumorvol = 55,  
  diagnosis = "neoplasia"), interval = "prediction")
```

	fit	lwr	upr
1	905.7322	-456.1596	2267.624

```
predict(model2, newdata = data_frame(tumorvol = 55,  
  diagnosis = "von H-L"), interval = "prediction")
```

	fit	lwr	upr
1	1299.003	-23.21001	2621.215

The broom package

```
tidy(model2)
```

	term	estimate	std.error
1	(Intercept)	417.2040	317.98858
2	log(tumorvol)	220.0463	93.57335
3	diagnosisneoplasia	-893.3017	658.56474
4	log(tumorvol):diagnosisneoplasia	124.7791	154.53279

	statistic	p.value
1	1.3120092	0.19857086
2	2.3515913	0.02481442
3	-1.3564372	0.18416834
4	0.8074602	0.42518351

The broom package

```
newdat <- augment(model2)
```

```
head(newdat, 2)
```

```
  p.ne log.tumorvol. diagnosis  .fitted  .se.fit  
1  289         2.564949   von H-L 981.6115 131.6002  
2  294         3.465736 neoplasia 718.9758 238.3433  
  .resid      .hat    .sigma    .cooksd .std.resid  
1 -692.6115 0.04312915 631.2164 0.01406830 -1.1173571  
2 -424.9758 0.14146982 638.3783 0.02158137 -0.7237946
```


The broom package

```
glance(model2)
```

```
   r.squared adj.r.squared   sigma statistic
1 0.2903844    0.2258739 633.6815   4.501351
      p.value df    logLik      AIC      BIC deviance
1 0.009374592  4 -289.0914 588.1827 596.2373 13251224
  df.residual
1           33
```

The Western Collaborative Group Study

Original description: 3524 men aged 39-59 and employed in the San Francisco Bay or Los Angeles areas were enrolled in 1960 and 1961. In addition to determinations of behavior pattern, the initial examination included medical and parental history, socioeconomic factors, exercise, diet, smoking, alcohol consumption, diet, serum lipid and lipoprotein studies, blood coagulation studies, and cardiovascular examination.

- <http://www.epi.umn.edu/cvdepi/study-synopsis/western-collaborative-group-study/>

The WCGS data describe 3,154 subjects and 22 variables. For now, let's examine a few interesting variables, and a sample of 500 of the observations.

```
wcgs.full <- read.csv("wcgs.csv")  
wcgs.full <- tbl_df(wcgs.full)
```

Select the variables of interest, and sample 500 subjects

```
set.seed(43101)
wcgs1 <-
  wcgs.full %>%
  select(id, age, chol, arcus, dibpat, bmi,
         wghtcat, smoke, ncigs, chd69) %>%
  sample_n(500, replace = FALSE)
```

Resulting tibble

```
# A tibble: 500 x 10
```

	id	age	chol	arcus	dibpat	bmi	wghtcat
	<int>	<int>	<int>	<int>	<fctr>	<dbl>	<fctr>
1	6039	42	218	0	Type A	25.05680	140-170
2	3713	44	207	0	Type B	22.87093	170-200
3	3465	50	249	0	Type A	24.38980	140-170
4	3923	42	236	0	Type B	25.10714	170-200
5	3503	46	247	1	Type A	21.47629	140-170
6	13167	40	206	0	Type B	26.45372	170-200
7	11325	39	267	1	Type B	23.56510	140-170
8	3800	51	214	0	Type A	23.67245	140-170
9	12582	43	235	1	Type A	19.25676	140-170
10	19096	48	234	0	Type B	28.12608	170-200

```
# ... with 490 more rows, and 3 more variables:  
#   smoke <fctr>, ncigs <int>, chd69 <fctr>
```

Codebook

Name	Stored As	Type	Details (units, levels, etc.)
id	integer	(nominal)	ID #, nominal and uninteresting
age	integer	quantitative	age, in years - no decimal places
chol	integer	quantitative	total cholesterol, mg/dL
arcus	integer	(nominal)	arcus senilis present (1) or absent (0)
dibpat	factor (2)	(binary)	behavioral pattern: A or B
bmi	number	quantitative	body-mass index
wghtcat	factor (4)	(ordinal)	wt: < 140, 140-170, 170-200, > 200
smoke	factor (2)	(binary)	cigarette smoker: Yes or No
ncigs	integer	quantitative	number of cigarettes smoked per day
chd69	factor (2)	(binary)	CHD event: Yes or No

Summary of this sample (without id)

age		chol		arcus	
Min.	:39.00	Min.	:110.0	Min.	:0.0000
1st Qu.	:42.00	1st Qu.	:201.0	1st Qu.	:0.0000
Median	:45.00	Median	:224.0	Median	:0.0000
Mean	:46.19	Mean	:228.2	Mean	:0.2966
3rd Qu.	:50.00	3rd Qu.	:255.0	3rd Qu.	:1.0000
Max.	:59.00	Max.	:400.0	Max.	:1.0000
		NA's	:3	NA's	:1

dibpat		bmi		wghtcat		smoke	
Type A:	251	Min.	:17.22	< 140	: 34	No	:255
Type B:	249	1st Qu.	:22.87	> 200	: 44	Yes:	245
		Median	:24.39	140-170:	252		
		Mean	:24.49	170-200:	170		
		3rd Qu.	:25.84				
		Max.	:38.95				

ncigs		chd69	
Min.	: 0.00	No	:462
1st Qu.	: 0.00	Yes:	38
Median	: 0.00		
Mean	:11.62		
3rd Qu.	:20.00		
Max.	:70.00		

```
summary(wcgs1)[,-1]
```

A Key Research Question

Were the men with Type A behavioral patterns more likely to suffer a CHD event?

```
table(wcgs1$dibpat, wcgs1$chd69)
```

	No	Yes
Type A	222	29
Type B	240	9

What's not so great about this table?

Re-specify the factor to re-order of the levels

```
wcgs1$chdevent <-  
  factor(wcgs1$chd69, levels = c("Yes", "No"))  
tab1 <- table(wcgs1$dibpat, wcgs1$chdevent)  
tab1
```

	Yes	No
Type A	29	222
Type B	9	240

and this is **standard epidemiological format**.

Aha! A Two-by-Two Table!

```
twoby2(tab1) ## twoby2 is part of the Epi package
```

2 by 2 table analysis:

Outcome : Yes

Comparing : Type A vs. Type B

	Yes	No	P(Yes)	95% conf. interval
Type A	29	222	0.1155	0.0815 0.1613
Type B	9	240	0.0361	0.0189 0.0680

	95% conf. interval
Relative Risk:	3.1965 1.5450 6.6134
Sample Odds Ratio:	3.4835 1.6132 7.5221
Conditional MLE Odds Ratio:	3.4754 1.5589 8.5398
Probability difference:	0.0794 0.0334 0.1279

A 4 by 2 table

```
table(wcgs1$wghtcat, wcgs1$chdevent)
```

	Yes	No
< 140	1	33
> 200	3	41
140-170	20	232
170-200	14	156

Why is this a poor table?

An Improved 4 x 2 table

```
wcgs1$wghtcat <- factor(wcgs1$wghtcat,  
  levels = c("< 140", "140-170", "170-200", "> 200"))  
tab2 <- table(wcgs1$wghtcat, wcgs1$chdevent)  
knitr::kable(tab2)
```

	Yes	No
< 140	1	33
140-170	20	232
170-200	14	156
> 200	3	41

A Three-Way Table (Not Flattened)

```
table(wcgs1$dibpat, wcgs1$wghtcat, wcgs1$chdevent)
```

, , = Yes

	< 140	140-170	170-200	> 200
Type A	0	15	13	1
Type B	1	5	1	2

, , = No



	< 140	140-170	170-200	> 200
Type A	15	104	82	21
Type B	18	128	74	20

A Three-Way Table (Flattened)

```
fable(wcgs1$dibpat, wcgs1$wghtcat, wcgs1$chdevent)
```

		Yes	No
Type A	< 140	0	15
	140-170	15	104
	170-200	13	82
	> 200	1	21
Type B	< 140	1	18
	140-170	5	128
	170-200	1	74
	> 200	2	20

How can we describe the distribution of BMI?



Summary of BMI data


```
mosaic::favstats(wcgs1$bmi)
```

	min	Q1	median	Q3	max	mean
	17.22242	22.87091	24.3898	25.84016	38.94737	24.48819
	sd	n	missing			
	2.680243	500	0			

```
psych::describe(wcgs1$bmi)
```

	vars	n	mean	sd	median	trimmed	mad	min	max
X1	1	500	24.49	2.68	24.39	24.39	2.15	17.22	38.95
	range	skew	kurtosis	se					
X1	21.72	0.65		2.11	0.12				

Comparing BMI by Behavior Pattern



Numerical BMI summaries, by behavior pattern

```
by(wcgs1$bmi, wcgs1$dibpat, mosaic::favstats)
```

```
wcgs1$dibpat: Type A
```

min	Q1	median	Q3	max	mean
18.30729	23.09703	24.40514	26.02431	33.90239	24.64689
sd	n	missing			
2.700341	251	0			


```
-----  
wcgs1$dibpat: Type B
```

min	Q1	median	Q3	max	mean
17.22242	22.59412	24.27378	25.82449	38.94737	24.32823
sd	n	missing			
2.65565	249	0			

How about Total Cholesterol, instead?

Warning: Removed 3 rows containing non-finite values (stat_bin).

Warning: Removed 3 rows containing non-finite values (stat_qq).



Why did we get error warnings?

```
mosaic::favstats(wcgs1$chol)
```

min	Q1	median	Q3	max	mean	sd	n	missing
110	201	224	255	400	228.167	44.20081	497	3

```
psych::describe(wcgs1$chol)
```



	vars	n	mean	sd	median	trimmed	mad	min	max
X1	1	497	228.17	44.2	224	226.64	41.51	110	400

	range	skew	kurtosis	se
X1	290	0.33	0.36	1.98

What should we do about missing values?

- 1 `ggplot2` doesn't include missing values in the plot, but it does warn that they've been removed.
 - We could suppress that warning by setting `na.rm = TRUE` in the call to a geom like `geom_histogram` or `geom_qq`.

Plot with `na.rm = TRUE`



Classification of Missing Data

There are three classifications we will think about in 431. Subtleties abound, but these three will suffice for most practical work.

- MCAR: Missing *Completely at Random*. This is the desirable scenario for us. MCAR means that there is no relationship between the probability of a data point being missing, and any values in the data set, either missing or observed.
 - The missing data are just a random subset of the data.
 - This is one kind of “ignorable” missingness.

Classification of Missing Data

- MAR: Missing *at Random*, which is definitely an unfortunate name. Less desirable, but there is still some hope. MAR means that the probability of a data point being missing has nothing to do with the missing value that would have been observed, but does have something to do with the values of some other variable that you did observe.
 - The idea is that if we can control for this other variable in our analysis, then we can treat this missingness as just a random subset of the data after that adjustment, which will eventually be pretty straightforward.
 - This is another kind of “ignorable” missingness.

Classification of Missing Data

- MNAR: Missing *not at Random* is a more serious issue. Now, additional thought and some special methods may well be required. Here, there is a relationship between the probability that a value is missing and what the actual (missing) value is.
 - This is what we mean by “non-ignorable” missingness.
 - Multiple Imputation methods (and Maximum Likelihood approaches) assume the data are MAR or MCAR, so the important distinction, generally, is between “at random” vs. “not at random.”
- Some of these explanations come from [this link](#)

What should we do about missing values?

- ② At times you will want to try to understand what makes observations with missing values different from observations with meaningful recorded values, especially if we're thinking that the missing mechanism is MCAR or MAR.
 - We might, for instance, compare the BMI values or perhaps the smoking status for those with and without missing cholesterol values, using the `is.na()` function to make a new variable to indicate those subjects without a cholesterol level.
- Some of this material is drawn from R for Data Science

For Tuesday: Working with Tables

For Class 11, I'd like you to read a 1981 paper by A.S.C. Ehrenberg that you'll find linked on the Class 10 README.

The paper is called *The Problem of Numeracy* and it provides some very helpful tips for working with tables, in particular.

There are three key tips related to the development of tables, in practice, as described by Ehrenberg, and also by Howard Wainer¹ who concisely states them as:

- 1 Order the rows and columns in a way that makes sense.
- 2 Round - a lot!
- 3 ALL is different and important.

¹Visual Revelations (1997), Chapter 10.