

# 431 Class 10

Thomas E. Love

2017-09-28

# Today's Agenda

- The Western Collaborative Group Study
  - Building Tables Effectively
  - Dealing with Missingness
  - Scatterplot and Correlation Matrices

# Today's R Setup

```
library(Epi); library(viridis)
library(GGally); library(mice)
library(tidyverse)

source("Love-boost.R")
```

# The Western Collaborative Group Study

Original description: 3524 men aged 39-59 and employed in the San Francisco Bay or Los Angeles areas were enrolled in 1960 and 1961. In addition to determinations of behavior pattern, the initial examination included medical and parental history, socioeconomic factors, exercise, diet, smoking, alcohol consumption, diet, serum lipid and lipoprotein studies, blood coagulation studies, and cardiovascular examination.

- <http://www.epi.umn.edu/cvdepi/study-synopsis/western-collaborative-group-study/>

The WCGS data describe 3,154 subjects and 22 variables. For now, let's examine a few interesting variables, and a sample of 500 of the observations.

```
wcgs.full <- read.csv("wcgs.csv")  
wcgs.full <- tbl_df(wcgs.full)
```

# Select the variables of interest, and sample 500 subjects

```
set.seed(43101)
wcgs1 <-
  wcgs.full %>%
  select(id, age, chol, arcus, dibpat, bmi,
         wghtcat, smoke, ncigs, chd69) %>%
  sample_n(500, replace = FALSE)
```

# Resulting tibble

```
# A tibble: 500 x 10
```

	id	age	chol	arcus	dibpat	bmi	wghtcat
	<int>	<int>	<int>	<int>	<fctr>	<dbl>	<fctr>
1	6039	42	218	0	Type A	25.05680	140-170
2	3713	44	207	0	Type B	22.87093	170-200
3	3465	50	249	0	Type A	24.38980	140-170
4	3923	42	236	0	Type B	25.10714	170-200
5	3503	46	247	1	Type A	21.47629	140-170
6	13167	40	206	0	Type B	26.45372	170-200
7	11325	39	267	1	Type B	23.56510	140-170
8	3800	51	214	0	Type A	23.67245	140-170
9	12582	43	235	1	Type A	19.25676	140-170
10	19096	48	234	0	Type B	28.12608	170-200

```
# ... with 490 more rows, and 3 more variables:  
#   smoke <fctr>, ncigs <int>, chd69 <fctr>
```

# Codebook

Name	Stored As	Type	Details (units, levels, etc.)
id	integer	(nominal)	ID #, nominal and uninteresting
age	integer	quantitative	age, in years - no decimal places
chol	integer	quantitative	total cholesterol, mg/dL
arcus	integer	(nominal)	arcus senilis present (1) or absent (0)
dibpat	factor (2)	(binary)	behavioral pattern: A or B
bmi	number	quantitative	body-mass index
wghtcat	factor (4)	(ordinal)	wt: < 140, 140-170, 170-200, > 200
smoke	factor (2)	(binary)	cigarette smoker: Yes or No
ncigs	integer	quantitative	number of cigarettes smoked per day
chd69	factor (2)	(binary)	CHD event: Yes or No

# Summary of this sample (without id)

age		chol		arcus	
Min.	:39.00	Min.	:110.0	Min.	:0.0000
1st Qu.	:42.00	1st Qu.	:201.0	1st Qu.	:0.0000
Median	:45.00	Median	:224.0	Median	:0.0000
Mean	:46.19	Mean	:228.2	Mean	:0.2966
3rd Qu.	:50.00	3rd Qu.	:255.0	3rd Qu.	:1.0000
Max.	:59.00	Max.	:400.0	Max.	:1.0000
		NA's	:3	NA's	:1

dibpat		bmi		wghtcat		smoke	
Type A:	251	Min.	:17.22	< 140	: 34	No	:255
Type B:	249	1st Qu.	:22.87	> 200	: 44	Yes:	245
		Median	:24.39	140-170:	252		
		Mean	:24.49	170-200:	170		
		3rd Qu.	:25.84				
		Max.	:38.95				

ncigs		chd69	
Min.	: 0.00	No	:462
1st Qu.	: 0.00	Yes:	38
Median	: 0.00		
Mean	:11.62		
3rd Qu.	:20.00		
Max.	:70.00		

```
summary(wcgs1)[,-1]
```



# A Key Research Question

Were the men with Type A behavioral patterns more likely to suffer a CHD event?

```
table(wcgs1$dibpat, wcgs1$chd69)
```

	No	Yes
Type A	222	29
Type B	240	9

What's not so great about this table?

# Re-specify the factor to re-order of the levels

```
wcgs1$chdevent <-  
  factor(wcgs1$chd69, levels = c("Yes", "No"))  
tab1 <- table(wcgs1$dibpat, wcgs1$chdevent)  
tab1
```

	Yes	No
Type A	29	222
Type B	9	240

and this is **standard epidemiological format**.

# Aha! A Two-by-Two Table!

```
twoby2(tab1) ## twoby2 is part of the Epi package
```

2 by 2 table analysis:

-----  
Outcome : Yes

Comparing : Type A vs. Type B

	Yes	No	P(Yes)	95% conf. interval
Type A	29	222	0.1155	0.0815 0.1613
Type B	9	240	0.0361	0.0189 0.0680

	95% conf. interval
Relative Risk:	3.1965 1.5450 6.6134
Sample Odds Ratio:	3.4835 1.6132 7.5221
Conditional MLE Odds Ratio:	3.4754 1.5589 8.5398
Probability difference:	0.0794 0.0334 0.1279

## A 4 by 2 table

```
table(wcgs1$wghtcat, wcgs1$chdevent)
```

	Yes	No
< 140	1	33
> 200	3	41
140-170	20	232
170-200	14	156

Why is this a poor table?

# An Improved 4 x 2 table

```
wcgs1$wghtcat <- factor(wcgs1$wghtcat,  
  levels = c("< 140", "140-170", "170-200", "> 200"))  
tab2 <- table(wcgs1$wghtcat, wcgs1$chdevent)  
knitr::kable(tab2)
```

	Yes	No
< 140	1	33
140-170	20	232
170-200	14	156
> 200	3	41

# A Three-Way Table (Not Flattened)

```
table(wcgs1$dibpat, wcgs1$wghtcat, wcgs1$chdevent)
```

, , = Yes

	< 140	140-170	170-200	> 200
Type A	0	15	13	1
Type B	1	5	1	2

, , = No



	< 140	140-170	170-200	> 200
Type A	15	104	82	21
Type B	18	128	74	20

# A Three-Way Table (Flattened)

```
fable(wcgs1$dibpat, wcgs1$wghtcat, wcgs1$chdevent)
```

		Yes	No
Type A	< 140	0	15
	140-170	15	104
	170-200	13	82
	> 200	1	21
Type B	< 140	1	18
	140-170	5	128
	170-200	1	74
	> 200	2	20

# How can we describe the distribution of BMI?





# Summary of BMI data

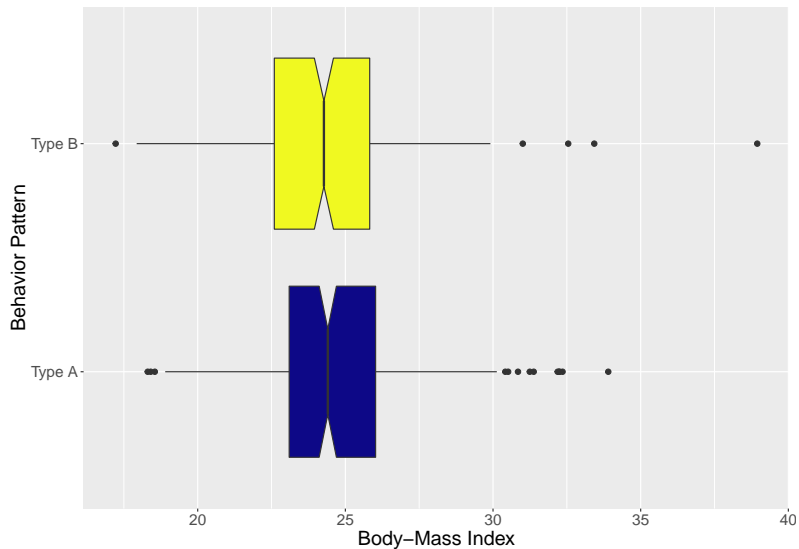
```
mosaic::favstats(wcgs1$bmi)
```

	min	Q1	median	Q3	max	mean
	17.22242	22.87091	24.3898	25.84016	38.94737	24.48819
	sd	n	missing			
	2.680243	500	0			

```
psych::describe(wcgs1$bmi)
```

	vars	n	mean	sd	median	trimmed	mad	min	max
X1	1	500	24.49	2.68	24.39	24.39	2.15	17.22	38.95
	range	skew	kurtosis	se					
X1	21.72	0.65		2.11	0.12				

# Comparing BMI by Behavior Pattern



# Numerical BMI summaries, by behavior pattern

```
by(wcgs1$bmi, wcgs1$dibpat, mosaic::favstats)
```

wcgs1\$dibpat: Type A

min	Q1	median	Q3	max	mean
18.30729	23.09703	24.40514	26.02431	33.90239	24.64689
sd	n	missing			
2.700341	251	0			


-----  
wcgs1\$dibpat: Type B

min	Q1	median	Q3	max	mean
17.22242	22.59412	24.27378	25.82449	38.94737	24.32823
sd	n	missing			
2.65565	249	0			

# How about Total Cholesterol, instead?

Warning: Removed 3 rows containing non-finite values  
(stat\_bin).

Warning: Removed 3 rows containing non-finite values  
(stat\_qq).



# Why did we get error warnings?

```
mosaic::favstats(wcgs1$chol)
```

min	Q1	median	Q3	max	mean	sd	n	missing
110	201	224	255	400	228.167	44.20081	497	3



```
psych::describe(wcgs1$chol)
```

	vars	n	mean	sd	median	trimmed	mad	min	max
X1	1	497	228.17	44.2	224	226.64	41.51	110	400
	range	skew	kurtosis	se					
X1	290	0.33	0.36	1.98					

# What should we do about missing values?

- 1 `ggplot2` doesn't include missing values in the plot, but it does warn that they've been removed.
  - We could suppress that warning by setting `na.rm = TRUE` in the call to a geom like `geom_histogram` or `geom_qq`.

# Plot with `na.rm = TRUE`



# Classification of Missing Data

There are three classifications we will think about in 431. Subtleties abound, but these three will suffice for most practical work.

- MCAR: Missing *Completely at Random*. This is the desirable scenario for us. MCAR means that there is no relationship between the probability of a data point being missing, and any values in the data set, either missing or observed.
  - The missing data are just a random subset of the data.
  - This is one kind of “ignorable” missingness.



# Classification of Missing Data

- MAR: Missing *at Random*, which is definitely an unfortunate name. Less desirable, but there is still some hope. MAR means that the probability of a data point being missing has nothing to do with the missing value that would have been observed, but does have something to do with the values of some other variable that you did observe.
  - The idea is that if we can control for this other variable in our analysis, then we can treat this missingness as just a random subset of the data after that adjustment, which will eventually be pretty straightforward.
  - This is another kind of “ignorable” missingness.

# Classification of Missing Data

- MNAR: Missing *not at Random* is a more serious issue. Now, additional thought and some special methods may well be required. Here, there is a relationship between the probability that a value is missing and what the actual (missing) value is.
  - This is what we mean by “non-ignorable” missingness.
  - Multiple Imputation methods (and Maximum Likelihood approaches) assume the data are at least MAR, so that’s the important distinction to make, generally.
- Some of these explanations come from [this link](#)

# What should we do about missing values?


- ② At times you will want to try to understand what makes observations with missing values different from observations with meaningful recorded values, especially if we're thinking that the missing mechanism is MCAR or MAR.
  - We might, for instance, compare the BMI values or perhaps the smoking status for those with and without missing cholesterol values, using the `is.na()` function to make a new variable to indicate those subjects without a cholesterol level.
- Some of this material is drawn from R for Data Science

# Do those missing cholesterol look unusual in terms of BMI?

First, we'll build a new (logical) variable (TRUE/FALSE) to indicate a missing cholesterol level, and then we'll plot the BMI distributions for each level of the new variable.

```
wcgs1 %>%  
  mutate(  
    nochol = is.na(chol)  
  ) %>%  
  ggplot(aes(x = bmi)) +  
  geom_freqpoly(aes(col = nochol), bins = 30) +  
  theme(text = element_text(size = 18))
```

# Do the BMIs of people without chol look different?



# Are the people without a cholesterol value unusual in terms of their smoking status?

```
temp1 <- table(is.na(wcgs1$chol), wcgs1$smoke)
knitr::kable(temp1)
```

	No	Yes
FALSE	253	244
TRUE	2	1

# What should we do about missing values?

## ③ How many missing values are there?

- If the missing values are more than, say, 5% of a variable, we're going to need some strong, almost heroic assumptions in order to feel confident about using such a variable in building a model or making an inference.
  - If the amount of missing data is very small relative to the size of the data as a whole, then leaving out a few samples and just running models or comparisons ignoring those observations may not be too damaging.
  - Depending on the situation, you may want to look for other fixes besides just dropping these cases and wiping out potentially useful data.
- 
- Some of this material comes from [this R-bloggers post](#)

# What should we do about missing values?

Could we **impute** missing values?

- One approach is *simple* imputation, where a single value is created to “fill in” the missing observation. This is pretty easy to do, but very rarely a good idea.
  - Rarely, substituting the mean is a reasonable thing to do, as it reduces variance in your estimate of the distribution, among other problems.
  - Sometimes, but still pretty rarely, substituting in a random value observed in the rest of the data set is a reasonable thing to do.
  - Better, although still problematic, imputation approaches use other variables in the data set to predict the missing value, and contain a random component. Using other variables preserves the relationships among variables in the imputations. The random component is important so that all missing values of a single variable are not all exactly equal. One example would be to use a regression equation to predict missing values, then add a random error term.
- See <http://www.theanalysisfactor.com/multiple-imputation-in-a-nutshell/>



# What's so bad about simple imputation?

Although there are several simple imputation approaches that solve many of the problems inherent in mean imputation, one problem remains. Because the imputed value is an estimate - a predicted value - there is uncertainty about its true value. Every statistic has uncertainty, measured by its standard error. Statistics computed using imputed data have even more uncertainty than its standard error measures. Your statistical package cannot distinguish between an imputed value and a real value.

Since the standard errors of statistics based on imputed values, such as sample means or regression coefficients, are too small, corresponding reported p-values are also too small. P-values that are reported as smaller than they, in reality, are, lead to Type I errors.

- It turns out that *multiple* imputation is a much better approach.
- Various types of “hot deck” procedures can help, too. See the `HotDeckImputation` package in R, or [this link](#)

# So what is multiple imputation?

Multiple imputation has solved this problem by incorporating the uncertainty inherent in imputation. It has four steps:

- 1 Create  $m$  sets of imputations for the missing values using an imputation process with a random component.
- 2 The result is  $m$  full data sets. Each data set will have slightly different values for the imputed data because of the random component.
- 3 Analyze each completed data set. Each set of parameter estimates will differ slightly because the data differs slightly.
- 4 Combine results, calculating the variation in parameter estimates.

# Multiple Imputation is amazing

Remarkably,  $m$ , the number of sufficient imputations, can be only 5 to 10 imputations, although it depends on the percentage of data that are missing. The result is unbiased parameter estimates and a full sample size, when done well.

Doing multiple imputation well, however, is not always quick or easy. First, it requires that the missing data be ignorable. Second, it requires a very good imputation model. Creating a good imputation model requires knowing your data very well and having variables that will predict missing values.

Source:

<http://www.theanalysisfactor.com/multiple-imputation-in-a-nutshell/>

# What will we do in 431?

- Often, we'll be willing to simply exclude the data with missing values from our graphs or other analyses.
- Sometimes, we'll be willing to assume (heroically) that the data are missing at random and we'll use a simple imputation approach, via the `mice` package.
- Later in the term (and definitely in 432) we'll move on up to multiple imputation, using `mice` sometimes and `Hmisc` at other times.

# Using mice to build imputations for chol

```
mice::md.pattern(wcgs1)
```

	id	age	dibpat	bmi	wghtcat	smoke	ncigs	chd69
496	1	1	1	1	1	1	1	1
3	1	1	1	1	1	1	1	1
1	1	1	1	1	1	1	1	1
	0	0	0	0	0	0	0	0

	chdevent	arcus	chol
496	1	1	1 0
3	1	1	0 1
1	1	0	1 1
	0	1	3 4

# Build 5 actual imputations using the “predictive mean matching” (pmm) approach

```
wcgs.temp <- mice(wcgs1,m=5,maxit=50,meth='pmm',seed=431)
```

iter	imp	variable
1	1	chol arcus
1	2	chol arcus
1	3	chol arcus
1	4	chol arcus
1	5	chol arcus
2	1	chol arcus
2	2	chol arcus
2	3	chol arcus
2	4	chol arcus
2	5	chol arcus
3	1	chol arcus

# View imputation results, summarized

```
> summary(wcgs.temp)
Multiply imputed data set
Call:
mice(data = wcgs1, m = 5, method = "pmm", maxit = 50, seed = 431)
Number of multiple imputations: 5
Missing cells per column:
      id      age      chol      arcus      dibpat      bmi      wghtcat      smoke      ncigs      chd69      chdevent
      0         0         3         1         0         0         0         0         0         0         0

Imputation methods:
      id      age      chol      arcus      dibpat      bmi      wghtcat      smoke      ncigs      chd69      chdevent
      "pmm"    "pmm"    "pmm"    "pmm"    "pmm"    "pmm"    "pmm"    "pmm"    "pmm"    "pmm"    "pmm"

VisitSequence:
 chol arcus
  3     4

PredictorMatrix:
      id age chol arcus dibpat bmi wghtcat smoke ncigs chd69 chdevent
id      0  0  0     0     0     0  0     0  0     0  0     0
age      0  0  0     0     0     0  0     0  0     0  0     0
chol     1  1  0     1     1     1  1     1  1     1  1     0
arcus     1  1  1     0     1     1  1     1  1     1  1     0
dibpat    0  0  0     0     0     0  0     0  0     0  0     0
bmi       0  0  0     0     0     0  0     0  0     0  0     0
wghtcat   0  0  0     0     0     0  0     0  0     0  0     0
smoke     0  0  0     0     0     0  0     0  0     0  0     0
ncigs     0  0  0     0     0     0  0     0  0     0  0     0
chd69     0  0  0     0     0     0  0     0  0     0  0     0
chdevent  0  0  0     0     0     0  0     0  0     0  0     0

Random generator seed value: 431
```

# Inspect the imputed values, if you like

```
wcgs.temp$imp$chol
```

	1	2	3	4	5
62	137	268	211	269	224
237	212	244	295	258	282
472	255	205	252	205	178

```
wcgs.temp$imp$arcus
```

	1	2	3	4	5
480	0	1	0	0	1



## Simple Imputation: Complete data with, let's say, the fourth of the five imputations we built

```
completedwcgs1 <- mice::complete(wcgs.temp,4)
```

### favstats with and without imputation

```
mosaic::favstats(wcgs1$chol)
```

min	Q1	median	Q3	max	mean	sd	n	missing
110	201	224	255	400	228.167	44.20081	497	3

```
mosaic::favstats(completedwcgs1$chol)
```

min	Q1	median	Q3	max	mean	sd	n	missing
110	201	224	255	400	228.262	44.13794	500	0

# Build a Linear Model without imputation

```
> modelFit0 <- with(wcgs1,lm(chol ~ bmi * dibpat))  
> summary(modelFit0)
```

```
Call:  
lm(formula = chol ~ bmi * dibpat)
```

Residuals:

Min	1Q	Median	3Q	Max
-117.432	-28.537	-4.019	25.859	175.580

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	185.096	25.576	7.237	1.77e-12	***
bmi	1.974	1.031	1.914	0.0562	.
dibpatType B	70.345	36.237	1.941	0.0528	.
bmi:dibpatType B	-3.325	1.471	-2.261	0.0242	*

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 43.74 on 493 degrees of freedom

(3 observations deleted due to missingness)

Multiple R-squared: 0.02664, Adjusted R-squared: 0.02072

F-statistic: 4.498 on 3 and 493 DF, p-value: 0.003984

```
> confint(modelFit0)
```

	2.5 %	97.5 %
(Intercept)	134.84532681	235.3476621
bmi	-0.05226551	4.0000876
dibpatType B	-0.85177283	141.5423274
bmi:dibpatType B	-6.21478658	-0.4358619

# Multiple imputation and pooling

Suppose that the next step in our analysis is to fit a linear model to the data. You may ask what imputed data set to choose. The `mice` package makes it again very easy to fit a a model to each of the imputed data sets and then pool the results together

```
> modelFit1 <- with(wcgs.temp, lm(chol ~ bmi * dibpat))  
> round(summary(pool(modelFit1)), 3)
```

	est	se	t	df	Pr(> t )	lo 95	hi 95	nmis	fmi	lambda
(Intercept)	185.409	25.776	7.193	434.927	0.000	134.749	236.069	NA	0.034	0.029
bmi	1.958	1.039	1.884	438.572	0.060	-0.084	4.000	0	0.032	0.028
dibpat2	71.446	36.278	1.969	479.010	0.049	0.163	142.730	NA	0.016	0.012
bmi:dibpat2	-3.362	1.472	-2.283	480.131	0.023	-6.255	-0.469	NA	0.016	0.012

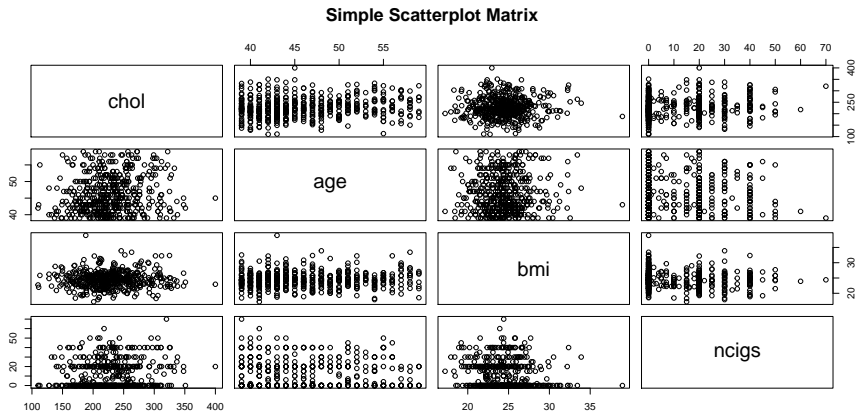
## Details of linear model after pooling

`modelFit1` contains the results of the fitting performed over the imputed data sets, while the `pool()` function pools them all together.

- `fmi` = fraction of missing information
- `lambda` = proportion of total variance attributable to the missing data
- Note that if we were looking at a strict alpha of 0.05, we'd have a significant `dibpat2` main effect now, when we didn't before.

# Multivariable Descriptions: A Scatterplot Matrix

```
pairs (~ chol + age + bmi + ncigs,  
       data=wcgs1, main="Simple Scatterplot Matrix")
```




# Correlation Matrix (after imputation)

```
completedwcgs1 <- mice::complete(wcgs.temp,1)  
round(cor(completedwcgs1[c("chol", "age", "bmi")]),3)
```

	chol	age	bmi
chol	1.000	0.133	0.036
age	0.133	1.000	0.029
bmi	0.036	0.029	1.000

# Using GGally for a Correlation Matrix

```
tempdat <- completedwgs1 %>%  
  select(chol, age, bmi, ncigs)  
  
ggcorr(tempdat, name = "Pearson r", label = TRUE)
```




# My Favorite Scatterplot Matrix

My favorite way to augment this plot adds loess smooths to the upper panel, and correlations in the lower panel, with histograms down the diagonal. To do this, we first create two functions (these modifications come from Chang's R Graphics Cookbook), called `panel.hist` and `panel.cor`.

These functions are in the `Love-boost.R` script.

# Augmented Scatterplot Matrix





# Code for Augmented Scatterplot Matrix


```
pairs (~ chol + age + bmi + ncigs, data=completedwcgs1,  
      main="Augmented Scatterplot Matrix",  
      upper.panel = panel.smooth,  
      diag.panel = panel.hist,  
      lower.panel = panel.cor)
```

# Using GGally for a Scatterplot Matrix (Code)

```
tempdat <- completedwchs1 %>%  
  select(chol, age, bmi, ncigs)  
  
ggpairs(tempdat, title = "Scatterplot Matrix via ggpairs")
```

# Using GGally for a Scatterplot Matrix (Result)

Scatterplot Matrix via ggpairs



## For Tuesday: Working with Tables

For Class 11, I'd like you to read a 1981 paper by A.S.C. Ehrenberg that you'll find linked on the Class 10 README.

The paper is called *The Problem of Numeracy* and it provides some very helpful tips for working with tables, in particular.

There are three key tips related to the development of tables, in practice, as described by Ehrenberg, and also by Howard Wainer<sup>1</sup> who concisely states them as:

- 1 Order the rows and columns in a way that makes sense.
- 2 Round - a lot!
- 3 ALL is different and important.

---

<sup>1</sup>Visual Revelations (1997), Chapter 10.