

# 431 Class 12

Thomas E. Love

2017-10-05

# Today's Agenda

- ① Some Thoughts on dplyr and its verbs
- ② The Printer Case Study
- ③ Dealing with Missing Data via Imputation
- ④ Setting up the first Quiz

# Setting Up Quiz 1

There are a total of 41 questions, 18 worth 2 points, 18 worth 3 points, 4 worth 2.5 points, and 1 that affirms your work is yours alone.

- Please select or type in your best response for each question. The questions are not arranged in any particular order, and you should answer all of them.
- You must complete this quiz by Noon on Monday, 2016-10-09. You will have the opportunity to edit your responses after completing the quiz, but this must be completed by the deadline.
- If you wish to complete part of the quiz and then return to it later, please scroll down to the end of the quiz and complete the **affirmation** (Question 41). The affirmation is required, and you will have to complete it in order to exit the quiz and save your progress. You will then be presented with a link to “Edit your progress” which you will want to bookmark, so you can return to it easily.

# Quiz 1: Main item types.

Fake Quiz is at <https://goo.gl/forms/hw37w3BrpibPDGQ03>

- ① Short Answer Questions
  - ② Multiple Choice
  - ③ Checkboxes
  - ④ Matching
- 
- You are welcome to consult the materials provided on the course website, but you are not allowed to discuss the questions on this quiz with anyone other than Professor Love or the Teaching Assistants, who may be reached at 431-help at case dot edu.

# Fake Quiz: Question A

## Fake Quiz for Demonstration Purposes

This is a FAKE quiz. NOT the REAL Quiz. Among other things, this FAKE quiz has only 4 items. The real one has 41.

Your email address (**tel3@case.edu**) will be recorded when you submit this form. Not you? [Switch account](#)

\* Required

### Fake Question A

Which of the statements below is true about outliers? (Check all that apply.)

- ☐ Outliers are values with Z scores below 2.
- ☐ Outliers indicate that something may be wrong with the data collection process.
- ☐ Outliers aren't important and should be identified and then ignored.
- ☐ None of these statements are true.

# Fake Quiz: Question B

## Fake Question B

Match the description of a relationship to a likely Pearson correlation coefficient.

	$r = 0$	$r = -0.3$	$r = 0.7$	$r = -0.7$	$r = 1$
A linear model fits the data very well, but not perfectly, and has a negative slope.	<input type="radio"/>	<input type="radio"/>	<input type="radio"/>	<input type="radio"/>	<input type="radio"/>
A loess smooth looks like a straight line with a negative slope, but the points are extremely widely scattered around the line, with a lot of variation shown.	<input type="radio"/>	<input type="radio"/>	<input type="radio"/>	<input type="radio"/>	<input type="radio"/>
Using <code>geom_smooth(method = "lm")</code> produces a horizontal line.	<input type="radio"/>	<input type="radio"/>	<input type="radio"/>	<input type="radio"/>	<input type="radio"/>

# Fake Quiz: Question C

## Fake Question C

What percentage of the observations drawn from a Normal distribution with mean 100 and variance 100 will be in the range of 80 to 120?

- ☐ Less than 20%
- ☐ 20 - 39%
- ☐ 40 - 59%
- ☐ 60 - 79%
- ☐ 80% or more

# Fake Quiz: Affirmation

## Affirmation Question \*

Please type in your name to indicate that you have not consulted with anyone else about this quiz except for Dr. Love and the teaching assistants, and that your answers are yours and yours alone. Just type in your full name.

Your answer \_\_\_\_\_

A copy of your responses will be emailed to tel3@case.edu.

SUBMIT



## Fake Quiz for Demonstration Purposes

Your response has been recorded.

[Edit your response](#)

# Today's R Setup

```
library(mice); library(tidyverse)  
  
source("Love-boost.R")
```

# dplyr basics: The Key Verbs

Six key functions:

- Pick observations by their values (`filter()`).
- Reorder the rows (`arrange()`).
- Pick variables by their names (`select()`).
- Collapse many values down to a single summary (`summarise()`).
- Create new variables with functions of existing variables (`mutate()`).
- Change the scope of another function from operating on the whole data set to operating on it group-by-group (`group_by()`)

*All of this comes from Wickham and Grolemund, R for Data Science, Chapter 5*

<http://r4ds.had.co.nz/transform.html#introduction-2>

# dplyr basics: How the verbs work

- The first argument is a data frame (or tibble).
- The second arguments describe what to do with the data frame. You can refer to columns in the data frame directly without using \$.
- The result is a new data frame.

We'll work with the `wcgs` data.

```
wcgs <- read.csv("wcgs.csv") %>% tbl_df(wcgs)
```

```
# A tibble: 3,154 x 22
```

	id	age	agec	height	weight	lnwght	wghtcat
	<int>	<int>	<fctr>	<int>	<int>	<dbl>	<fctr>
1	2343	50	46-50	67	200	5.298317	170-200
2	3656	51	51-55	73	192	5.257495	170-200
3	3526	59	56-60	70	200	5.298317	170-200
4	22057	51	51-55	69	150	5.010635	140-170
5	12227	44	41-45	71	160	5.075174	140-170

# Filter rows with filter()

filter() allows you to subset observations based on their values.

```
wcgs.sub1 <- wcgs %>%  
  filter(dibpat == "Type A" & age > 49)  
wcgs.sub1
```

# A tibble: 522 x 22

	id	age	agec	height	weight	lnwght	wghtcat
	<int>	<int>	<fctr>	<int>	<int>	<dbl>	<fctr>
1	2343	50	46-50	67	200	5.298317	170-200
2	3656	51	51-55	73	192	5.257495	170-200
3	3526	59	56-60	70	200	5.298317	170-200
4	22057	51	51-55	69	150	5.010635	140-170
5	12681	50	46-50	71	195	5.273000	170-200
6	3284	59	56-60	72	206	5.327876	> 200
7	21071	54	51-55	67	152	5.023880	140-170
8	13371	55	51-55	72	185	5.220356	170-200


# Comparison and Logical Operators

Comparison Operator	Meaning
>	is greater than
>=	is greater than or equal to
<	is less than
<=	is less than or equal to
!=	is not equal to
==	is equal to

Logical (Boolean) Operator	Meaning
&	and
	or
!	not

Missing Values (NA in R) can make things tricky. They are contagious.  
Almost any operation involving an unknown value will also be unknown.

# The complete set of Boolean Operators



Source: <http://r4ds.had.co.nz/transform.html#logical-operators>

# Arrange rows with `arrange()`

`arrange()`, instead of selecting rows (like `filter()`), changes their order.

- Use `arrange(height)` to arrange in ascending order of height.  
Provide a second column name to break ties, if you like.
- Missing values are always sorted at the end.

```
wcgs %>%  
  arrange(desc(height), desc(weight))
```

```
# A tibble: 3,154 x 22
```

	id	age	agec	height	weight	lnwght	wghtcat
	<int>	<int>	<fctr>	<int>	<int>	<dbl>	<fctr>
1	12012	47	46-50	78	250	5.521461	> 200
2	2145	41	41-45	78	220	5.393628	> 200
3	12680	43	41-45	78	190	5.247024	170-200
4	13512	42	41-45	77	220	5.393628	> 200
5	12620	49	46-50	77	210	5.347107	> 200
6	11209	45	41-45	77	195	5.273000	170-200



# Select columns with select()

select() lets you zoom in on the columns you actually want to use based on the names of the variables. R for Data Science lays out some helper functions within select() for use in bigger data sets.

```
wcgs.sub2 <- wcgs %>%  
  select(id, age, height, weight, dibpat, smoke, behpat)  
wcgs.sub2
```

```
# A tibble: 3,154 x 7
```

	id	age	height	weight	dibpat	smoke	behpatt
	<int>	<int>	<int>	<int>	<fctr>	<fctr>	<fctr>
1	2343	50	67	200	Type A	Yes	A1
2	3656	51	73	192	Type A	Yes	A1
3	3526	59	70	200	Type A	No	A1
4	22057	51	69	150	Type A	No	A1
5	12927	44	71	160	Type A	No	A1
6	16029	47	64	158	Type A	Yes	A1

# Grouped summaries with summarize()

summarise() or summarize() collapses a data frame to a single row.

```
wcgs.sub2 %>%  
  summarize(mean.ht = mean(height, na.rm=TRUE),  
            sd.ht = sd(height, na.rm=TRUE)) %>%  
  round(digits = 2)
```

```
# A tibble: 1 x 2  
  mean.ht sd.ht  
    <dbl> <dbl>  
1   69.78   2.53
```

# Using the pipe (%>%) to filter and summarize

```
wcgs.sub2 %>%  
  filter(dibpat == "Type A") %>%  
  summarize(pearson.r = cor(height, weight),  
    spearman.r = cor(height, weight, method = "spearman")) %>%  
  round(digits = 3) %>%  
  knitr::kable()
```

pearson.r	spearman.r
0.534	0.542

# Using `group_by()` with `summarize` to look group-by-group

```
wcgs.sub2 %>%  
  group_by(behpat) %>%  
  summarize(  
    pearson.r = round(cor(height, weight), 3) ) %>%  
  knitr::kable()
```

behpat	pearson.r
A1	0.571
A2	0.526
B3	0.524
B4	0.557

## Using `group_by()` to look at separated groups

You might have tried this approach instead, but it throws an error...

```
wcgs.sub2 %>%  
  group_by(behpat) %>%  
  summarize(  
    pearson.r = cor(height, weight)) %>%  
  round(digits = 3) %>%  
  knitr::kable()
```

- Why doesn't this work?

## Using `group_by()` to look at separated groups

You might have tried this approach instead, but it throws an error...

```
wcgs.sub2 %>%  
  group_by(behpat) %>%  
  summarize(  
    pearson.r = cor(height, weight)) %>%  
  round(digits = 3) %>%  
  knitr::kable()
```

- Why doesn't this work?
- When R sees the `round` command, it tries to apply it to every element of the table, including the behavior pattern labels, which aren't numbers. So it throws an error.

# Add new variables with mutate()

mutate() adds new columns that are functions of existing columns to the end of your data set.

Suppose we want to calculate the weight/height ratio for each subject.

```
wcgs.sub3 <- wcgs.sub2 %>%  
  mutate(wh.ratio = weight / height)  
wcgs.sub3
```

```
# A tibble: 3,154 x 8
```

	id	age	height	weight	dibpat	smoke	behpat
	<int>	<int>	<int>	<int>	<fctr>	<fctr>	<fctr>
1	2343	50	67	200	Type A	Yes	A1
2	3656	51	73	192	Type A	Yes	A1
3	3526	59	70	200	Type A	No	A1
4	22057	51	69	150	Type A	No	A1
5	12927	44	71	160	Type A	No	A1
6	16029	47	64	158	Type A	Yes	A1

# On Coding and dplyr

- ① Learn dplyr, and use it to do most of your data management within R.
  - dplyr is mostly about these key verbs, and piping, for our purposes
  - some tasks produce results which be confusing, we're here to help
- ② dplyr is most useful in combination with other elements of the tidyverse, most prominently ggplot2.
- ③ Hmisc doesn't play nicely with dplyr, so don't load the whole Hmisc library, just call individual functions you need with, for example, `Hmisc::describe` or `Hmisc::smean.cl.boot`



# The Printer Case, Setup

## The Printer Case

Your firm is located in a five-story building<sup>1</sup>. Each floor has its own printer/copier in a copy room. The firm owns these machines but must pay for paper, toner and occasional maintenance. Each employee has a key that opens the copy room door on their floor only and does not have access to machines on other floors. Because the printer/copiers are “free goods” right now, you suspect that the firm’s printing costs could be cut drastically. To test this, you performed an experiment. The third and fifth floors were chosen because these two floors have had about the same usage rates in the past. Each person on the fifth floor was given a card to operate the fifth floor machine. These employees were told that their card would generate a daily accounting of their printer activity. Fifth floor employees have also been told that they will not be *charged* for their use of the machine, but they certainly know that *someone* will have some sense of individual usage patterns. To establish a basis of comparison, the group on the third floor has not been converted to the card system. The third floor machine has an internal mechanism that totals the number of copies made each day, but you do not know *who* is doing *what*, and the third floor employees have no reason to believe that they are being monitored.

# The Printer Case, Main Table

You collected data from the machines over the last 50 working days. The data are in the table below and can be downloaded from the web in the **printer.csv** file. There are three variables: **DAY**, which indicates the day; **FIFTH**, the number of copies made on the 5<sup>th</sup> floor; and **THIRD**, the number made on the 3<sup>rd</sup> floor.

Will the card accounting system effectively lower usage if implemented across the firm?

Day	Fifth	Third	Day	Fifth	Third	Day	Fifth	Third	Day	Fifth	Third
1	750	340	14	570	370	27	390	270	39	270	400
2	710	540	15	570	720	28	420	670	40	250	130
3	700	210	16	560	670	29	380	660	41	210	440
4	720	530	17	500	460	30	370	240	42	240	130
5	690	550	18	480	320	31	370	500	43	190	250
6	670	350	19	550	370	32	360	480	44	160	330
7	660	590	20	510	570	33	350	560	45	130	300
8	640	520	21	520	120	34	330	310	46	120	110
9	670	360	22	460	190	35	280	390	47	180	740
10	620	420	23	470	710	36	300	610	48	150	700
11	580	160	24	440	620	37	310	690	49	110	150
12	590	470	25	400	180	38	290	410	50	100	580
13	610	380	26	410	640						

# The Printer Case Discussion, Part 1

Fifty days of data:

- Fifth floor employees were given a card to operate their printer.
- Third floor employees were not.

- 1 Is this a randomized trial or an observational study?
- 2 What is the outcome we are studying?
- 3 What are the two treatments/exposures/interventions being compared?
- 4 What controls are in place as part of the study's design?
- 5 **Key Question:** Will the card accounting system effectively lower usage if implemented across the firm?

# The Printer Case Discussion


Go.

# Printer Case: Numerical Summary

```
printer <- read.csv("printer.csv") %>% tbl_df(  
summary(printer)
```

Day	Fifth	Third
Min. : 1.00	Min. :100.0	Min. :110.0
1st Qu.:13.25	1st Qu.:282.5	1st Qu.:302.5
Median :25.50	Median :415.0	Median :415.0
Mean :25.50	Mean :426.2	Mean :428.2
3rd Qu.:37.75	3rd Qu.:577.5	3rd Qu.:577.5
Max. :50.00	Max. :750.0	Max. :740.0

# Printer Case: Scatterplot ( $r = 0.11$ )



# Printer Case: Gather the Columns


First, we'll gather up the data so that we can plot it more easily.

```
printer2 <- tidyr::gather(printer, Floor, Pages, -Day)
printer2
```

```
# A tibble: 100 x 3
```

	Day	Floor	Pages
	<int>	<chr>	<int>
1	1	Fifth	750
2	2	Fifth	710
3	3	Fifth	700
4	4	Fifth	720
5	5	Fifth	690
6	6	Fifth	670
7	7	Fifth	660
8	8	Fifth	640
9	9	Fifth	670


# Printer Case: Comparison Boxplot



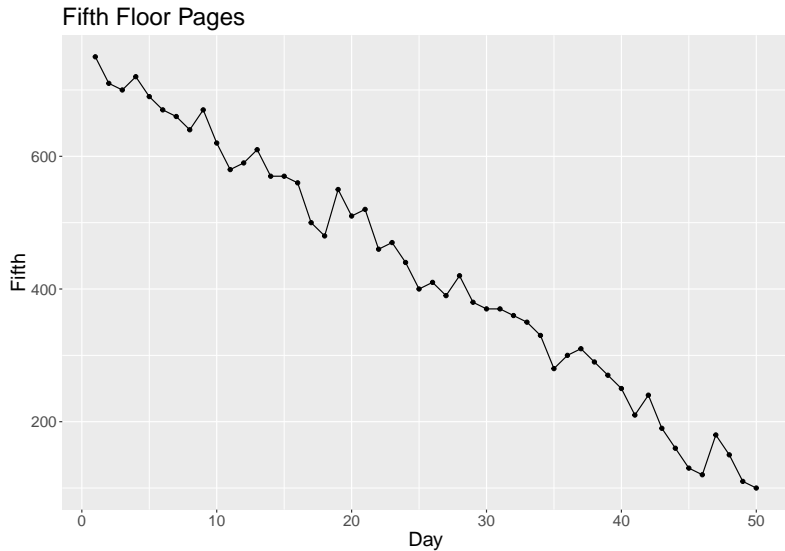


# Printer Case: Third Floor

Third Floor Pages

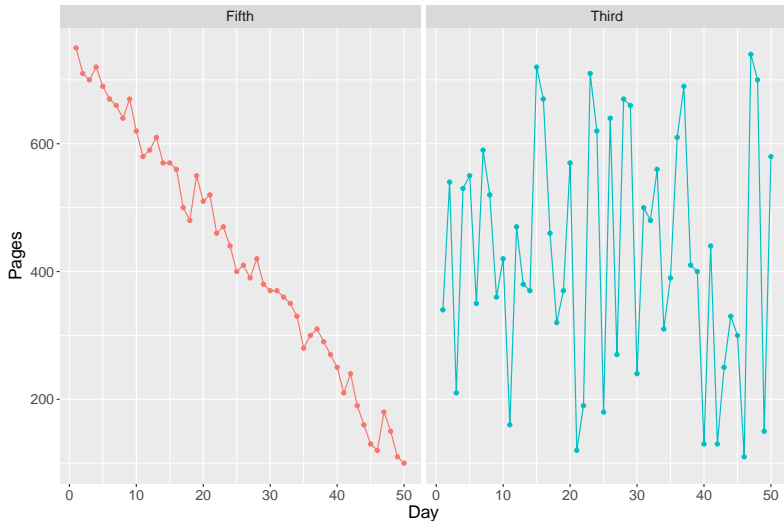


# Printer Case: Fifth Floor



# Comparing the Patterns over Time

Monitoring on Fifth Floor Reduced Pages



# Back to WCGS: Select variables, sample 500 subjects

```
set.seed(43101)
wcgs1 <-
  wcgs %>%
  select(id, age, chol, arcus, dibpat, bmi,
         wghtcat, smoke, ncigs, chd69) %>%
  sample_n(500, replace = FALSE)
```

# Dealing with Missing Data

## ① Missing Completely at Random (MCAR)

- The missing data are just a random subset of the data.

## ② Missing at Random (MAR)

- MAR means that the probability of a data point being missing has nothing to do with the missing value that would have been observed, but does have something to do with the values of some other variable that you did observe.
- Multiple Imputation assumes the missingness is MCAR or MAR.

## ③ Missing not at Random (“non-ignorable” missingness)

- Here, there is a relationship between the probability that a value is missing and what the actual (missing) value is.

# What should we do about missing values?


- ① `ggplot2` doesn't include missing values in the plot, but it does warn that they've been removed.
- ② At times you will want to try to understand what makes observations with missing values different from observations with meaningful recorded values, especially if we're thinking that the missing mechanism is MCAR or MAR.
  - We might, for instance, compare the BMI values or perhaps the smoking status for those with and without missing cholesterol values, using the `is.na()` function to make a new variable to indicate those subjects without a cholesterol level.

# Do NA cholesterol look unusual in terms of BMI?

This code builds a new (logical) variable (TRUE/FALSE) to indicate a missing cholesterol level, and then we'll plot the BMI distributions for each level of the new variable.

```
wcgs1 %>%  
  mutate(  
    nochol = is.na(chol)  
  ) %>%  
  ggplot(aes(x = bmi)) +  
  geom_freqpoly(aes(col = nochol), bins = 30) +  
  theme(text = element_text(size = 18))
```

# Do the BMIs of people without chol look different?





# Are the people without a cholesterol value unusual in terms of their smoking status?

```
temp1 <- table(is.na(wcgs1$chol), wcgs1$smoke)
knitr::kable(temp1)
```

	No	Yes
FALSE	253	244
TRUE	2	1

# What should we do about missing values?

## ③ How many missing values are there?

- If the missing values are more than, say, 5% of a variable, we're going to need some strong, almost heroic assumptions in order to feel confident about using such a variable in building a model or making an inference.
  - If the amount of missing data is very small relative to the size of the data as a whole, then leaving out a few samples and just running models or comparisons ignoring those observations may not be too damaging.
  - Depending on the situation, you may want to look for other fixes besides just dropping these cases and wiping out potentially useful data.
- 
- Some of this material comes from [this R-bloggers post](#)

# What should we do about missing values?

Could we **impute** missing values?

- One approach is *simple* imputation, where a single value is created to “fill in” the missing observation. This is pretty easy to do, but very rarely a good idea.
  - Rarely, substituting the mean is a reasonable thing to do, as it reduces variance in your estimate of the distribution, among other problems.
  - Sometimes, but still pretty rarely, substituting in a random value observed in the rest of the data set is a reasonable thing to do.
  - Better, although still problematic, imputation approaches use other variables in the data set to predict the missing value, and contain a random component. Using other variables preserves the relationships among variables in the imputations. The random component is important so that all missing values of a single variable are not all exactly equal. One example would be to use a regression equation to predict missing values, then add a random error term.
- See <http://www.theanalysisfactor.com/multiple-imputation-in-a-nutshell/>

# What's so bad about simple imputation?

Although there are several simple imputation approaches that solve many of the problems inherent in mean imputation, one problem remains. Because the imputed value is an estimate - a predicted value - there is uncertainty about its true value. Every statistic has uncertainty, measured by its standard error. Statistics computed using imputed data have even more uncertainty than its standard error measures. Your statistical package cannot distinguish between an imputed value and a real value.

Since the standard errors of statistics based on imputed values, such as sample means or regression coefficients, are too small, corresponding reported p-values are also too small. P-values that are reported as smaller than they, in reality, are, lead to Type I errors.

- It turns out that *multiple* imputation is a much better approach.
- Various types of “hot deck” procedures can help, too. See the HotDeckImputation package in R, or [this link](#)

# So what is multiple imputation?

Multiple imputation has solved this problem by incorporating the uncertainty inherent in imputation. It has four steps:

- 1 Create  $m$  sets of imputations for the missing values using an imputation process with a random component.
- 2 The result is  $m$  full data sets. Each data set will have slightly different values for the imputed data because of the random component.
- 3 Analyze each completed data set. Each set of parameter estimates will differ slightly because the data differs slightly.
- 4 Combine results, calculating the variation in parameter estimates.

# Multiple Imputation is amazing

Remarkably,  $m$ , the number of sufficient imputations, can be only 5 to 10 imputations, although it depends on the percentage of data that are missing. The result is unbiased parameter estimates and a full sample size, when done well.

Doing multiple imputation well, however, is not always quick or easy. First, it requires that the missing data be ignorable. Second, it requires a very good imputation model. Creating a good imputation model requires knowing your data very well and having variables that will predict missing values.

Source:

<http://www.theanalysisfactor.com/multiple-imputation-in-a-nutshell/>

# What will we do in 431?

- Often, we'll be willing to simply exclude the data with missing values from our graphs or other analyses.
- Sometimes, we'll be willing to assume (heroically) that the data are missing at random and we'll use a simple imputation approach, via the `mice` package.
- Later in the term (and definitely in 432) we'll move on up to multiple imputation, using `mice` sometimes and `Hmisc` at other times.

# Using mice to build imputations for chol

```
md.pattern(wcgs1)
```

	id	age	dibpat	bmi	wghtcat	smoke	ncigs	chd69	arcus
496	1	1	1	1	1	1	1	1	1
3	1	1	1	1	1	1	1	1	1
1	1	1	1	1	1	1	1	1	0
	0	0	0	0	0	0	0	0	1

	chol
496	1 0
3	0 1
1	1 1
	3 4



# Build 5 actual imputations using the “predictive mean matching” (pmm) approach

```
wcgs.temp <- mice(wcgs1,m=5,maxit=50,meth='pmm',seed=431)
```

iter	imp	variable
1	1	chol arcus
1	2	chol arcus
1	3	chol arcus
1	4	chol arcus
1	5	chol arcus
2	1	chol arcus
2	2	chol arcus
2	3	chol arcus
2	4	chol arcus
2	5	chol arcus
3	1	chol arcus

# View imputation results, summarized

```
> summary(wcgs.temp)
Multiply imputed data set
Call:
mice(data = wcgs1, m = 5, method = "pmm", maxit = 50, seed = 431)
Number of multiple imputations: 5
Missing cells per column:
      id      age      chol      arcus      dibpat      bmi      wghtcat      smoke      ncigs      chd69      chdevent
      0         0         3         1         0         0         0         0         0         0         0
Imputation methods:
      id      age      chol      arcus      dibpat      bmi      wghtcat      smoke      ncigs      chd69      chdevent
      "pmm"     "pmm"     "pmm"     "pmm"     "pmm"     "pmm"     "pmm"     "pmm"     "pmm"     "pmm"     "pmm"
VisitSequence:
 chol arcus
   3     4
PredictorMatrix:
      id age chol arcus dibpat bmi wghtcat smoke ncigs chd69 chdevent
id      0  0  0      0      0  0  0      0  0  0      0      0
age      0  0  0      0      0  0  0      0  0  0      0      0
chol      1  1  0      1      1  1      1  1  1      1      0
arcus      1  1  1      0      1  1      1  1  1      1      0
dibpat      0  0  0      0      0  0      0  0  0      0      0
bmi        0  0  0      0      0  0      0  0  0      0      0
wghtcat     0  0  0      0      0  0      0  0  0      0      0
smoke       0  0  0      0      0  0      0  0  0      0      0
ncigs       0  0  0      0      0  0      0  0  0      0      0
chd69       0  0  0      0      0  0      0  0  0      0      0
chdevent    0  0  0      0      0  0      0  0  0      0      0
Random generator seed value: 431
```

# Inspect the imputed values, if you like

```
wcgs.temp$imp$chol
```

	1	2	3	4	5
62	128	224	194	178	234
237	235	245	202	198	226
472	208	227	210	242	212

```
wcgs.temp$imp$arcus
```

	1	2	3	4	5
480	0	0	0	0	0

## Simple Imputation: Complete data with, let's say, the fourth of the five imputations we built

```
completedwcgs <- mice::complete(wcgs.temp,4)
```

### favstats with and without imputation

```
mosaic::favstats(wcgs1$chol)
```

min	Q1	median	Q3	max	mean	sd	n	missing
110	201	224	255	400	228.167	44.20081	497	3

```
mosaic::favstats(completedwcgs$chol)
```

min	Q1	median	Q3	max	mean	sd	n
110	200.75	224	255	400	228.034	44.14974	500
missing							
							0

# Build a Linear Model without imputation

```
> modelFit0 <- with(wcgs1,lm(chol ~ bmi * dibpat))
> summary(modelFit0)
```

Call:  
lm(formula = chol ~ bmi \* dibpat)

Residuals:

Min	1Q	Median	3Q	Max
-117.432	-28.537	-4.019	25.859	175.580

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	185.096	25.576	7.237	1.77e-12 ***
bmi	1.974	1.031	1.914	0.0562 .
dibpatType B	70.345	36.237	1.941	0.0528 .
bmi:dibpatType B	-3.325	1.471	-2.261	0.0242 *

---  
Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 43.74 on 493 degrees of freedom  
(3 observations deleted due to missingness)  
Multiple R-squared: 0.02664, Adjusted R-squared: 0.02072  
F-statistic: 4.498 on 3 and 493 DF, p-value: 0.003984

```
> confint(modelFit0)
```

	2.5 %	97.5 %
(Intercept)	134.84532681	235.3476621
bmi	-0.05226551	4.0000876
dibpatType B	-0.85177283	141.5423274
bmi:dibpatType B	-6.21478658	-0.4358619

# Multiple imputation and pooling

Suppose that the next step in our analysis is to fit a linear model to the data. You may ask what imputed data set to choose. The `mice` package makes it again very easy to fit a a model to each of the imputed data sets and then pool the results together

```
> modelFit1 <- with(wcgs.temp, lm(chol ~ bmi * dibpat))
> round(summary(pool(modelFit1)), 3)
```

	est	se	t	df	Pr(> t )	lo 95	hi 95	nmis	fmi	lambda
(Intercept)	185.409	25.776	7.193	434.927	0.000	134.749	236.069	NA	0.034	0.029
bmi	1.958	1.039	1.884	438.572	0.060	-0.084	4.000	0	0.032	0.028
dibpat2	71.446	36.278	1.969	479.010	0.049	0.163	142.730	NA	0.016	0.012
bmi:dibpat2	-3.362	1.472	-2.283	480.131	0.023	-6.255	-0.469	NA	0.016	0.012

## Details of linear model after pooling

`modelFit1` contains the results of the fitting performed over the imputed data sets, while the `pool()` function pools them all together.

- `fmi` = fraction of missing information
- `lambda` = proportion of total variance attributable to the missing data
- Note that if we were looking at a strict alpha of 0.05, we'd have a significant `dibpat2` main effect now, when we didn't before.

# Link to the Quiz

will be provided by 3 PM Thursday 2017-10-05.